

Key Concepts

Contents

- LLMs
- Generation
- LLMResult

LLMs

Wrappers around Large Language Models (in particular, the “generate” ability of large language models) are at the core of LangChain functionality. The core method that these classes expose is a `generate` method, which takes in a list of strings and returns an LLMResult (which contains outputs for all input strings). Read more about LLMResult.

This interface operates over a list of strings because often the lists of strings can be batched to the LLM provider, providing speed and efficiency gains. For convenience, this class also exposes a simpler, more user friendly interface (via `__call__`). The interface for this takes in a single string, and returns a single string.

Generation

The output of a single generation. Currently in LangChain this is just the generated text, although could be extended in the future to contain log probs or the like.

LLMResult

The full output of a call to the `generate` method of the LLM class. Since the `generate` method takes as input a list of strings, this returns a list of results. Each result consists of a list of generations (since you can request N generations per input string). This class contains a `llm_output` attribute which contains provider-specific information about the call.